Bengali Classification and Evidence-Bounded Evaluation

Two benchmark tasks | two pretrained encoders | three seeds per task-encoder pair

BENCHMARK CORPORA

Fake News: 94,883 articles
BD-SHS: 50,281 comments

P1: ORIGINAL PARTITIONS

Archived source-label evaluation
2 tasks x 2 encoders x 3 seeds
12 original fine-tuning runs
Fixed source-labelled test sets

P2 V1: EXACT-OVERLAP-CONTROLLED

NFKC + casefold + punctuation filtering
Remove matching earlier-split rows
2 tasks x 2 encoders x 3 seeds
12 separate cleaned-split runs

PRIMARY PERFORMANCE

Seed-wise accuracy and macro-F1
Paired significance and class-wise metrics
Original P1 test labels only

CLEAN-SPLIT SENSITIVITY

12 seed-wise P2 accuracies
Zero normalized exact-text split matches
Residual near matches: not leakage-free

SUPPORTING VALIDITY AUDITS AND REPORTING

Human-label review (360 news + 240 comments) | selected SHAP / LIME / IG probes
Auxiliary 12-article sliding-window inference (not P1 test accuracy)
Scope-qualified conclusions; no held-out-source or fully leakage-free claim