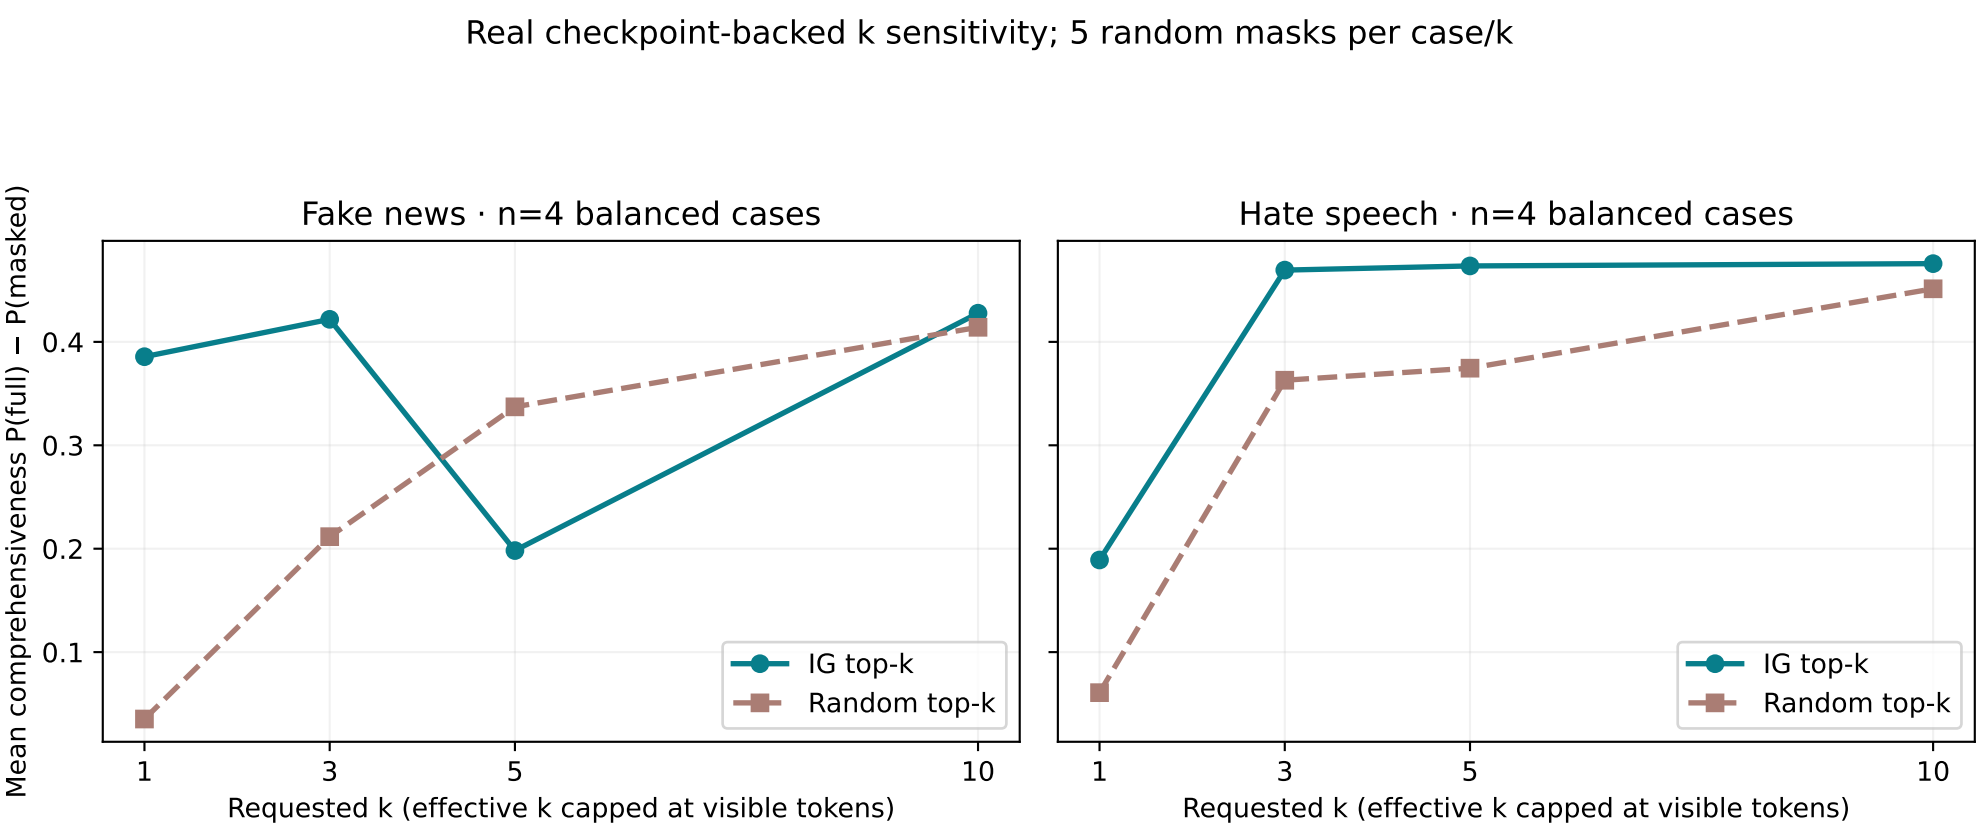
Real checkpoint-backed k sensitivity; 5 random masks per case/k
Fake news · n=4 balanced cases
Hate speech · n=4 balanced cases
Mean comprehensiveness P(full) − P(masked)
0.4
0.3
0.2
0.1
IG top-k
Random top-k
1
3
5
10
Requested k (effective k capped at visible tokens)